

# RAG System Explained - Simple Version


## What is this system?


This is a smart chatbot that helps employees find information from company documents. Instead of searching through hundreds of pages, you just ask a question and get an answer with the source.

## Overall Architecture Diagram

```
USER asks question
↓
SYSTEM searches documents
↓
SYSTEM finds best matches
↓
SYSTEM ranks them by relevance
↓
AI generates answer from matches
↓
USER gets answer + source
```

### Visual Flow:





## End-to-End Data Flow

---

1. Load Documents
  - System reads PDF, Word, HTML files
  - Example: HR\_Policies/vacation.pdf
2. Break into Chunks
  - Split each document into small pieces (1500 characters each)
  - Why? So we can find exact relevant parts
  - Example: One chunk = "Vacation Policy: Employees get 15 days..."
3. Convert to Numbers (Embeddings)
  - Similar chunks get similar numbers
4. Store in Database
  - Save all chunks with their numbers in FAISS database
  - Also save: which file, which category

## Part 2: Answering Questions

**What happens:** System finds relevant info and generates answer.

```

USER: "What is the vacation policy?"
↓
STEP 1: Convert question to numbers (same way as documents)
↓
STEP 2: Search database
  - Find 10 chunks with most similar numbers
  - Filter by category if detected (e.g., HR)
↓
STEP 3: Rerank the 10 chunks
  - Score each one more carefully
  - Keep only top 3 most relevant
↓
STEP 4: Generate answer
  - AI reads the 3 chunks
  - Writes answer using ONLY info from those chunks
  - Adds source citation
↓
ANSWER: "Employees receive 15 days of vacation per year..."
SOURCE: HR_Policies/vacation.pdf

```

## How Embedding Works

---

**Simple Explanation:** Embedding = Converting text into numbers so computers can compare them.

**Example:**

- Text: "vacation policy"
- Embedding: [0.23, -0.45, 0.67, ... 1536 numbers total]

**Why it's useful:**

- Similar meanings → Similar numbers
- "vacation policy" and "time off rules" get similar numbers
- System can find related info even if exact words don't match

**In our system:**

1. We use AWS Titan Embeddings model
2. Every document chunk gets converted numbers

3. Every user question also gets converted numbers
4. We compare these numbers to find matches

#### Visual:

Document: "Employees get 15 vacation days"

↓ Titan Embeddings

Numbers: [0.12, 0.34, -0.56, ...]

Question: "How many vacation days?"

↓ Titan Embeddings

Numbers: [0.15, 0.32, -0.54, ...]

Compare → Very similar! → This document matches the question

## How Retrieval Works

---

**Simple Explanation:** Retrieval = Finding the most relevant document chunks for a question.

#### Step-by-Step:

##### 1. Detect Category (Optional)

- System looks for keywords in question
- "vacation" → HR\_Policies category
- "API" → Engineering\_Docs category
- Helps narrow down search

##### 2. Convert Question to Numbers

- Question: "What is vacation policy?"
- Becomes: [0.15, 0.32, -0.54, ... 1536 numbers]

##### 3. Search Database

- Compare question numbers with all document chunk numbers
- Use FAISS (fast search tool)
- Find 10 chunks with closest numbers

##### 4. Return Results

- Get 10 most similar chunks

- o Each has: text content, source file, category

## Example:

Question: "What is the vacation policy?"

Search Results (Top 10):

1. "Employees receive 15 days vacation..." (HR\_Policies/vacation.pdf)
2. "Vacation days can be carried over..." (HR\_Policies/vacation.pdf)
3. "Time off requests must be submitted..." (HR\_Policies/timeoff.pdf)
- ... (7 more)

## Where Reranking Happens

---

Simple Explanation: Reranking = Double-checking which chunks are REALLY the most relevant.

### Why we need it:

- First search is fast but not perfect
- Reranking is slower but more accurate
- Better to spend extra time and get the right answer

When it happens: After retrieval, before generation.

### How it works:

1. We have 10 chunks from retrieval
2. Re-embed each chunk
  - o Convert each chunk to numbers again (fresh calculation)
  - o More accurate than stored embeddings
3. Calculate similarity scores
  - o Compare question numbers with each chunk numbers
  - o Use dot product (multiply and add)
  - o Higher score = more relevant
4. Sort by score
  - o Rank chunks from highest to lowest score

## 5. Keep top 3

- Only send best 3 chunks to AI
- Reduces noise, improves answer quality

Example:

After Retrieval (10 chunks):

Chunk 1: Score 0.85  
Chunk 2: Score 0.92 ← Best  
Chunk 3: Score 0.78  
Chunk 4: Score 0.88 ← 2nd best  
Chunk 5: Score 0.65  
Chunk 6: Score 0.81 ← 3rd best  
Chunk 7: Score 0.55  
Chunk 8: Score 0.70  
Chunk 9: Score 0.60  
Chunk 10: Score 0.50

After Reranking (keep top 3):

- ✓ Chunk 2 (0.92)
- ✓ Chunk 4 (0.88)
- ✓ Chunk 6 (0.81)

# How Generation Works

---

**Simple Explanation:** Generation = AI reads the relevant chunks and writes an answer in plain English.

**Step-by-Step:**

## 1. Prepare Context

- Take the 3 best chunks from reranking
- Format them nicely for AI
- Add labels (category/filename)

## 2. Create Prompt

- Combine: chunks + question + instructions
- Instructions tell AI: "Only use info from these chunks"

## 3. Send to AI (Claude 3 Haiku)

- AI reads everything
- Writes answer based ONLY on the chunks
- Cannot make up information

#### 4. Add Citation

- System adds source file at the end
- Shows which document the answer came from

#### 5. Return to User

- User sees answer + source

Example:

Input to Chatbot:

Context:  
[HR\_Policies/vacation.pdf]

"Employees receive 15 days of vacation per year. Vacation days can be carried over to the next year up to a maximum of 5 days."

[HR\_Policies/vacation.pdf]  
"Vacation requests must be submitted at least 2 weeks in advance through the HR portal."

Question: "What is the vacation policy?"

Instructions: Answer using ONLY the context above.

Output from Chatbot:

Employees receive 15 days of vacation per year. Vacation days can be carried over to the next year, but only up to 5 days maximum. To request vacation, you must submit your request at least 2 weeks in advance through the HR portal.

---

 Source Referenced: HR\_Policies/vacation.pdf

## Complete Example: User Question to Answer

User asks: "How many vacation days do I get?"

### Step 1: Embedding (0.2 seconds)

- Convert question to numbers: [0.15, 0.32, -0.54, ...]

### Step 2: Retrieval (0.05 seconds)

- Search database for similar chunks
- Find 10 candidates from HR\_Policies

### Step 3: Reranking (2 seconds)

- Score all 10 chunks carefully
- Keep best 3:
  - i. "Employees receive 15 days vacation..."
  - ii. "Vacation days can be carried over..."
  - iii. "New employees receive prorated vacation..."

### Step 4: Generation (2.5 seconds)

- AI reads the 3 chunks
- Writes: "Employees receive 15 days of vacation per year..."
- Adds source: HR\_Policies/vacation.pdf

Total time: ~5 seconds

User sees:

Employees receive 15 days of vacation per year. Vacation days can be carried over to the next year up to a maximum of 5 days.

---

 Source Referenced: [HR\\_Policies/vacation.pdf](#)

## Summary

---

What the system does:

1. Stores all company documents as searchable numbers
2. When you ask a question, finds relevant parts

3. Ranks them by relevance
4. AI reads them and writes an answer
5. Shows you the source

### Why it's useful:

- Fast: Get answers in 5 seconds
- Accurate: Only uses real documents
- Traceable: Always shows source
- Smart: Understands meaning, not just keywords

### Technology used:

- AWS Titan: Converts text to numbers
- FAISS: Fast search database
- Claude AI: Writes answers
- Python: Connects everything